

COMPLETE READER EDITION

Liang Wenfeng

Complete Four-Hour Investor Meeting Transcript

Unofficial public transcript / English reader edition

DURATION

Approx. 4 hours

MEETING

Investor discussion

LANGUAGE

English

COMPLETE TEXT

19,775 source words / 69 speaker turns / no transcript cuts

CALM TYPOGRAPHY / CLICKABLE NAVIGATION

DEEPSEEK / INVESTOR MEETING / UNOFFICIAL PUBLIC COPY

READER NOTE

Built for a long, demanding conversation

A four-hour transcript is easier to understand when the page helps you keep your place. This edition preserves the supplied text while adding a clickable contents page, thematic chapter breaks, distinct speaker markers and repeated continuation labels inside long monologues.

- 01** Use the contents page or PDF bookmarks to jump between the opening address and the four Q&A blocks.
- 02** Speaker colors remain consistent: teal for Liang Wenfeng, blue for investors and coral for the host.
- 03** Long turns repeat the speaker label periodically, so a page break never leaves you wondering who is speaking.
- 04** Editorial chapter titles and glossary notes are clearly separated from the transcript itself.

EDITORIAL STATUS

This edition preserves the complete user-provided transcript. Chapter titles, glossary definitions and navigation notes are editorial additions and are not part of the meeting record.

19,775

SOURCE WORDS

69

SPEAKER TURNS

15

OPENING
TAKEAWAYS

100%

COMPLETE
ENGLISH TEXT



CONTENTS

Navigate the conversation

Every entry is clickable. PDF bookmarks provide the same navigation in the side panel of most readers.


—	Key takeaways	4
01	Opening address	6
02	Strategy and industry outlook	32
03	Continual learning, talent and compute	43
04	Scaling, self-improvement and embodiment	56
05	Products, data and commercialization	61
06	TileLang and closing discussion	66
—	Glossary	69
—	Source and editorial status	70

OPENING OVERVIEW

Key takeaways

The source file labels this overview as lightly edited. The statements below are preserved from the supplied text and should not be treated as verified quotations.

- 01** NVIDIA's ecosystem is not irreplaceable; meaningful cracks could appear within a year. Domestic alternatives like Huawei are expected to reach functional parity within a year — even a 1-2× price premium would be acceptable.
- 02** Anyone trying to extract windfall profits from AI will most likely fail. Those who practice restraint in capturing value are the ones more likely to win. OpenAI's current predicament is fundamentally tied to its “winner-take-all” mindset.
- 03** The so-called “Singularity” is not a sudden eruption — it is a gradual evolutionary process, and the sequence matters: first solve continuous learning, then achieve self-iteration, and only after that, embodied intelligence.
- 04** The AI technology roadmap, broadly speaking: Language Models → Chain of Thought (CoT) → Agents → Continuous Learning → Self-Iteration (Intelligence Singularity) → Embodied Intelligence. We are currently in the mid-stage of the Agent phase. Continuous learning is the next critical threshold that must be crossed.
- 05** “Restraint” is the core strategy: deliberately opting out of the fight for consumer traffic, not building a super app. Near-term monetization serves only as a safety net; all core resources are focused entirely on AGI.
- 06** The China-US AI gap is roughly one year, but China has used only about 1/20th of the other side's compute to get there. Talent is not the problem; the core bottleneck remains compute (chips).

- 
- 07 At this stage, data matters most. The bottleneck for high-quality data is not money — it is time. In fact, roughly half of our core researchers are currently involved in data labeling.
 - 08 Even in the worst-case scenario where technological progress stalls, API services alone would be sufficient to sustain a publicly listed company. Monetization is not an existential concern.
 - 09 Organizationally, we do not overemphasize overtime, KPIs, or complex hierarchies. We rely more on vision to drive people. Research requires a certain “looseness” — excessive pressure actually reduces efficiency.
 - 10 Domestic chips will continue to face capacity constraints over the next 1-3 years, but on a 5-year horizon, the problem is largely solvable.
 - 11 There are too many companies in China building foundation models. The field will eventually converge to 3-4 players, and the industry as a whole will not settle into a high-margin structure.
 - 12 Domains like video generation, 3D, and world models are not closely tied to the ceiling of intelligence. They are more about commercial showcase value, so we do not treat them as core investment areas.
 - 13 As long as the core team remains stable and incentives are in place, AGI is highly likely to be achieved. Everything else is just a matter of timing — maybe a year earlier or later.
 - 14 China's long-term competitive advantage in AI lies in “systemic low cost,” much like manufacturing. It is a structural advantage.
 - 15 Once continuous learning is cracked, AI will be able to participate in its own research and iteration, and general intelligence will enter an accelerated phase of development.

CHAPTER 01

Opening address

Vision, restraint, open source and the commercial logic behind DeepSeek's founding choices.

LIANG WENFENG

FOUNDER

Welcome, everyone.

When we first started this company, our original intention was never about how much money we'd eventually make, or going to the capital markets, or going public — none of that was part of the plan. We simply didn't have that intention.

The first few dozen people never thought about it that way. If they had, they wouldn't have joined in the first place.

So in general, we embarked on this endeavor with tremendous goodwill toward the world. We genuinely believe this is something beneficial to humanity — something that matters beyond money.

Of course, later on, once it became clear how enormous the stakes were, other temptations emerged. That's a separate story. But our founding intention, our vision, and the vision we've maintained to this day — none of it was designed around maximizing commercial returns. I think this point is fairly critical.

About twenty years ago, the person I admired most in management was Jack Welch, the former CEO of GE. Looking back now, most of what he said may no longer hold true, but he got one thing absolutely right: the most important thing for a company is its vision. Managing a large company doesn't rest on your rules and regulations — it rests on vision.

What is vision? Vision is not a slogan you hang on the wall. Vision is how you act, not what you say. It's how you actually operate. I've forgotten Welch's exact words, but that was the gist of it.

So how do we manage so many people? How did we organize? The truth is, we don't really have an organization — we are vision-driven. We organize around a shared vision. We don't really have a formal structure.

This has its upsides and its downsides. Going forward, we'll figure out how to play to the strengths and mitigate the weaknesses, but this is our defining characteristic. We don't operate on the premise of “I need to achieve some KPI” or “there are performance reviews.” There is only the vision.

LIANG WENFENG · CONTINUED

FOUNDER

And this vision isn't even written down. It's not codified anywhere. Nothing has ever been written out. The vision lives in how we do things and in our attitude toward the world.

Perhaps every single person in the company understands the vision a little differently. Their individual visions may vary somewhat. But on the big-picture direction, there is alignment. I think we still carry tremendous goodwill toward the world and just want to do something meaningful. That's what binds us together.

Let me lay things out first, and then we can open it up to questions. I'll try to frame everything that follows around this vision.

This vision is genuine. It's not fabricated. We genuinely think this way and genuinely act this way. Otherwise, you simply couldn't explain many of the things we do.

Why are we so committed to open source? Because the vision itself demands open source. Without this vision, you couldn't bring people together in this way.

For example, Company X also open-sources, but their open source is different from ours. Theirs carries a sense of being forced — they don't feel it's their true intention. For us, this is exactly our intention.

And on the topic of open source, we were very clear about it from the very beginning. First, there's the vision. Second, we believe that for AI to succeed commercially, open source is actually advantageous. This sounds a

bit contradictory, a bit counterintuitive, because historically, open source and commercialization have been in tension.

But I think AI is different from what came before. In the past, a software company's market might amount to a few tens of billions of dollars. If you went open source, that would be gone — you might be left with only a few hundred million or a few billion. But AI is big enough that it could ultimately account for, say, 10% of global GDP. That is an enormous number.

LIANG WENFENG · CONTINUED

FOUNDER

No single party can monopolize this. You cannot keep it all to yourself — you absolutely have to share it with others, or you simply won't survive. This is fundamentally different from open-sourcing a piece of software in the past, because those software markets were never that big. But AI is just too big.

If we tried to monopolize the benefits, history would inevitably leave us behind. I think this above all is an objective law, a historical perspective. It's not that if I don't open-source, I get to monopolize the market. That doesn't hold up theoretically — it simply doesn't align with objective reality. You would encounter enormous resistance, and there would always be other forces preventing you from achieving that goal.

Under these circumstances, I don't think you need to adhere rigidly to traditional business thinking. You need a mechanism that ensures the benefits you capture for yourself are bounded — only then can you succeed.

Restraint is necessary. I believe restraint is necessary. If we want to pull off AI in our hands, the very first thing is restraint. You can't go around thinking that some percentage of global GDP should all be mine, or that some percentage of China's GDP should all be mine. The more you think that way, the less likely you are to succeed.

So from the beginning, we felt restraint was essential. The more restrained you are, the more likely you are to pull this off. That's a commercial consideration — a macro-level one, of course. I think it's intuitive, or at least it aligns with my intuition. Or at minimum, I genuinely think this way.

We don't have all that many other advantages. We don't have special capabilities. We're not richer than others, and our people aren't necessarily better than those at other companies. Honestly, we aren't.

Think about it: two years ago when we founded this company, we didn't have much money, we didn't have many GPUs, we had no name recognition, no real convening power. We were just a group of thoroughly ordinary people. We really were just a group of ordinary people.

If there's a narrative I like, it's that a group of ordinary people achieved something extraordinary — not that a group of geniuses achieved something extraordinary. This is deeply connected to our restraint, deeply tied to our vision. They're all of a piece.

Now, does open source conflict with commercialization? I think in the domain of AI, if you're not restrained, you won't even survive. Open source is part of that restraint. And our restraint doesn't only manifest in open source — it shows up in many other areas as well.

LIANG WENFENG · CONTINUED

FOUNDER

But overall, we don't have to agonize over open source, and we don't have to agonize over the question of restraint. The more restrained you are, the more likely you are to succeed. At least so far, the evidence bears this out — so far, it holds up as an explanation.

Otherwise, there's no way to explain how we managed to succeed. We didn't have any special weapons. Our starting point was very low, our resources were very few, and our people were essentially a random collection of ordinary individuals. I myself am just a university graduate — not even from a top-tier school.

This restraint is also part of our vision. AI is too big, and the stakes are too large. We are very restrained. As long as we can pull this off, the ultimate benefits will be enormous. Take just a sliver, and it's already huge. So right now, there's absolutely no need to think about which piece of the pie to take or how to take it. I genuinely believe there's no need to consider that at all, because the pie itself is big enough. Take just a tiny fraction, and it's more than enough.

That's why we've said before: we only aim to earn a reasonable profit. It's about your intention, not about how large the profit could be — those are different things.

This is not some abstract philosophy — it's reflected in our API pricing. For API pricing, we consider a reasonable profit to be roughly recovering the cost of a server in ten months. That feels reasonable to us. Given the current landscape — accounting for risk, upfront investment, and so on — if a server is amortized over three to five years from an accounting standpoint, but commercially we think recovering cost in about ten months is good enough, then OK. We're satisfied.

That's the logic behind our API pricing. Our V3.2 Flash and other models are all priced to recover hardware cost in ten months. That's our benchmark.

This is not profit maximization. If we were maximizing profit, we'd set prices higher. Because in this price range, user demand is inelastic — if I doubled the price, or raised it by even more, the consumption of tokens wouldn't change much. If I doubled the price, my total revenue would nearly double.

Hang on, let me double-check something... Ah, great.

Let me tell you a story. One of our models — at first, we were worried demand would be too high, so we initially priced it fairly high. The team wasn't happy about it. Later, I dropped the price to one-quarter of the original, and everyone was thrilled.

LIANG WENFENG · CONTINUED

FOUNDER

I think this is what we genuinely believe. Going back to the vision I mentioned earlier — we want this technology to be genuinely useful to people, not to maximize how much money we extract. We want to earn a reasonable profit while keeping it accessible for everyone.

I think this is how everyone else at the company feels too. When we cut prices, the company group chat was full of cheers. People were genuinely happy. Because this is the whole point — the reason we put so much effort

into making the model great: so it can be very inexpensive, very effective, and widely available for everyone to use to their heart's content.

That gives us joy. That's our motivation. That's our vision. That's the shared understanding that binds the company together and drives us forward. It's our internal consensus.

This is probably somewhat unusual. Because a price cut like that is certainly not good news for our competitors — they definitely aren't cheering. Your revenue, your ARR — cut it in half, and your ARR drops in half.

Right — this is one way we're different. We feel this is enough.

Internally, if I can recover cost in ten months, commercially I'm already very satisfied. Externally, we also feel this is a price point that makes people happy, that people are glad to see. It's win-win — the company and society, everyone wins together.

I think, OK — someone just commented on the screen that a ten-month payback means the profit is too high. And honestly, there is indeed still room for price reductions. There really is. On the model optimization side, there's also still room. So overall, the potential for further price reductions is still fairly significant.

But the thing about that ten-month payback cost: we can achieve it ourselves. Others can't. Companies like Alibaba or Tencent — they don't have our optimizations. Their costs are probably several times higher. There's still a lot of optimization work involved.

As for why we don't keep cutting prices — it's because there's no elasticity. If I cut prices further, demand won't increase more. Or if I cut further, the increase in demand becomes negligible. Because at this price, everyone can already afford it. Everyone feels the price is reasonable. No one is being priced out.

LIANG WENFENG · CONTINUED

FOUNDER

So further price cuts wouldn't bring the company more revenue, and they wouldn't bring society more value, because everyone is already satisfied

with the current price. Going even lower wouldn't meaningfully increase social welfare.

Right, OK. But on this point — when it comes to pricing, our starting point is definitely not maximizing company revenue or profit. That's not the basis for our decisions. This is part of our restraint. In the short term, higher prices might mean more revenue. But in the long term, it's far from clear.

Because I view restraint as a strategy. For me, restraint is a strategy. It's about sometimes giving something up to gain something else in return.

The same goes for open source. You can see it as a cost we bear, or you can see it as value we share. First, this value-sharing makes us happy internally — everyone is glad, employees feel a sense of achievement, and it strengthens our cohesion. And this value-sharing benefits society — society is happy, other industry players and ordinary people are happy.

So this restraint — the way I understand it — over the long run increases our probability of achieving AGI. When I consider any decision, I have zero doubt that AGI will carry enormous commercial value. On that basis, my priority isn't how to grab a bigger share, how to take more for myself. My priority is how to increase the probability that we succeed.

This restraint shows up in many other areas too. For example, during last year's Spring Festival, we suddenly got a massive influx of users. But we didn't chase after retaining those users, or try to monetize them, or rush to capture those commercial interests and cash in on the user base. We didn't fight for users, we didn't try to make money off them — but we worked very hard to serve them well.

We never entertained the thought of “I'm going to build the next super app, then compete with so-and-so, become the next ByteDance or the next Tencent.” Absolutely no such thinking. We could have done that, but we chose not to.

My understanding is that this, too, is part of restraint. Don't try to capture everything. You might feel like now that you have users, you could become the next ByteDance, so you go ahead and swallow that up. I think that could work commercially — it's a viable path. If last year we had thrown big

money at competing with ByteDance for users, that would have been one approach.

But we chose a very restrained path: I'm not going to fight you over this, because there's a watermelon waiting further down the road, and what's in front of us right now might just be sesame seeds. I shouldn't be grabbing at every sesame seed.

LIANG WENFENG · CONTINUED

FOUNDER

Granted, this sesame seed might be a big one. But I think compared to what's coming in AI, nothing up front really counts as big. Looking back now, our decision last year not to go all-out on the consumer side was probably the right call. Because we can now see there really are much bigger watermelons ahead. What was in front of us was indeed just some small sesame seeds.

If I had lots of money last year and built that consumer business up huge — what would I have actually gained? Nothing much.

These are my genuine thoughts. Because I think the AGI opportunity down the road is enormous. The AGI opportunity down the road will always be enormous. I don't even need to think about whether I'll have a seat at the table, or what my business model will look like at that point. We don't need to think about that at all. With an opportunity that large, you'll always find a way.

The sesame seeds up front — we'll pick them up too, but casually, in passing. We won't stop and make them the main focus. So last year's consumer DAU and all that — I think that was probably just a small thing. But we did pick it up. We maintained user engagement at relatively low cost, because it might be useful someday. Even though right now we don't know what the users are useful for — currently it's purely a cost center — in the future it might be valuable. Since it's something we could grab without much effort, we grabbed it.

Looking at this year, there's also a real possibility that our API or AI-related ARR revenue presents an opportunity. If demand can continue expanding, and if we can buy more GPUs, then ARR reaching several hundred million

dollars is quite plausible. If AI revenue can hit a billion dollars, then the company's cash flow could basically turn positive — covering our R&D expenses, covering all our costs.

So that's a possibility too. But we haven't made it a priority. We'll do it, but I don't think it's a top-priority matter. It's not our number-one concern, and it's not what we truly care about today.

The bigger opportunity is still ahead. The opportunities up front — last year's consumer side, this year's enterprise side — I think these are things we should do and do well, but they're not our goal. Or rather, most people in the company don't consider this to be all that important — not on the same level of importance as AGI.

Let me expand a bit on open source, since so many questions have touched on it. First, I believe we will continue to open-source, and our strongest models will likely also be open-sourced. Because I don't see any meaningful upside to keeping things closed. I just don't see it.

LIANG WENFENG · CONTINUED

FOUNDER

ByteDance keeps its models closed-source — what advantage does that give them? I can't see any.

Even if you open-source a model and tell everyone everything, the barrier to entry remains very high. For others to actually deploy and use it — that barrier is very high. Getting it to work at all is hard. On top of that, getting it to work at very low cost is extremely, extremely difficult. It's not easy at all.

It's not as if, the moment I open-source, someone else can trivially match my deployment costs. There's still a tremendous amount of work involved. Even though the principles are understood, not every company is willing — or has the will and capability — to organize the human resources needed to reach that level. I'm used to this. They may simply not be good at doing this kind of thing, because the internal resistance is too great. Controlling those costs is hard — there are many managerial and physical constraints.

This is also an advantage for a startup company. If you're too small, you lack the strength to pull it off. If you're a big company, you struggle to organize around it. Both ends have their challenges. So this falls into a sweet spot

unique to a company of our scale — our sweet point. If we were bigger, we might have other problems; if we were smaller, we wouldn't have enough strength.

So as for open source, I think we should approach it with the right pricing model. Right now, we should recognize that we're not forcing anyone. Because with our pricing model, I won't charge an exorbitant fee — I'll probably charge based on the same ten-month cost-recovery principle. Pricing at ten-month payback is enough to make independent third-party deployment unprofitable — they can't compete. A third party can't achieve that cost level. They simply can't.

So open source won't hurt our revenue. Of course, if I were aiming for a hundredfold profit, then open source would...

HOST

HOST

We can hear you, but it looks like your video dropped, boss. Maybe the call just transferred in.

LIANG WENFENG

FOUNDER

So, open source — I think it has absolutely no impact on our business model. The premise is that we only aim for a sixfold profit. Ten months to recover cost corresponds roughly to a sixfold profit. If we're only taking a sixfold profit, open source doesn't really affect anything.

But if you want a hundredfold profit, then open source does indeed get in the way — because a third party could deploy it at, say, twenty times the cost and undercut you.

Is this model sustainable over the long term? I think it is. Under our vision, I believe open source is sustainable over the long term, and that's what we intend to do.

You could say that restraint also enables you to thrive over the long haul. This strategy gives us more opportunities on the technology front and increases our probability of achieving AGI. We're more at ease. Think about it — we don't even need to work overtime, because it's just not that hard.

But for other players, it might be very hard, because they're trying to think about too many things at once. It's actually not that difficult at all. It really isn't.

From the outside, it may look like we chose a hard path — doing research, tackling the hardest problems — as if we're in hard mode. But in reality, we've given up a lot in other areas, which makes us still very effective and actually quite relaxed in how we work.

So my assessment on open source is that it's sustainable. There is no conflict between open source and commercial revenue, provided you're operating at a sixfold profit level. A sixfold profit may seem high, but it's really not. Given how efficient AI is right now, a reasonable profit is probably around that level. In the future, it might drop to, say, fourfold or threefold — I think that's... you can't really go lower than that. But it'll still be a significant profit.

Just from the perspective of selling API access alone — though I don't think selling API access is all that compelling — you can see that I don't perceive any conflict.

And I'm not at all worried about someone deploying our model and competing with us. Not worried in the slightest. In fact, we hope they do deploy it. We do our best to help the open-source community, to assist everyone in getting our models deployed. I'm not concerned they'll steal my business — the market is big enough. I'm only concerned they won't be able to deploy it properly, that some details won't be done right, and the performance will suffer, or their costs will end up high.

LIANG WENFENG · CONTINUED

FOUNDER

Right — there's no conflict here.

Last year, when the enterprise business started taking shape, I got a lot of questions along the lines of: if I open-source, won't my consumer business then conflict with it? Because I don't have a traffic advantage, and Tencent, say — with its vast user base — could deploy our open-source model and scoop up all the consumer users, taking them away from me.

But actually, it doesn't happen. There are many reasons for this.

Another question: is the open-source model we release the same as the model we deploy ourselves? Yes, it's the same. We wouldn't open-source an inferior model while keeping a better one for ourselves. That's not how we operate. It's the same model. This also demonstrates that there's really no conflict.

Over the entire past year, on the consumer side, we basically open-sourced everything — and we saw no conflict with our consumer services. None.

So that's the open-source piece.

Now, regarding the company's long-term vision — I think our goal should be AGI. Everyone may define AGI a little differently, but that doesn't stop us from treating AGI as our north star.

From a technical roadmap perspective, the path to AGI is actually fairly clear. With the current generation of AI technology, if you can articulate a problem very clearly and provide it with complete context and instructions, it already surpasses humans. But there's a caveat, a precondition: you have to give it complete context and complete instructions. And that precondition is very hard to fulfill.

Take today's meeting, for example. We're sitting here with a long, rich history of context — decades of context for many of us — that AI simply doesn't have. What AI can do right now is perform better than humans within a bounded context. But it still can't replace humans.

LIANG WENFENG · CONTINUED

FOUNDER

What's missing is continuous learning. Humans can learn continuously. You hire a new employee — they might take two months to get familiar with the company's environment and their work. After two months of onboarding, they're up and running. They can handle a lot of things. They can understand what you're saying — if you say, “Get Xiao Wang to come over,” they know who Xiao Wang is.

But with AI, because it lacks those two months of prior context, if you tell it “get Xiao Wang,” you have to spell out who Xiao Wang is, what his role is,

where he sits, how to find him, and what to watch out for when looking for him. You have to give AI all of that context. In that situation, AI can do the task, but you can't possibly feed it all the context — it's just not practical. So AI cannot replace your employee.

But if AI had the ability to learn continuously — just like your employee spending two months learning the ropes — then it could replace anyone. So what we're missing before the next leap is exactly this: learning to learn.

You can think of AI's development as a staircase. Last year's step was CoT — Chain of Thought. We discovered that through chain-of-thought reasoning, intelligence could reach a higher level. By having it think through problems on its own, the ceiling was raised and AI could do more things. So we climbed that step.

This year's step is Agents. We've found that with the Agent approach, even more tasks become feasible — the capability envelope expands, and the intelligence ceiling rises further.

Why staircase? Because every step is built on the foundation of the previous one. Agents depend on CoT, and CoT depends on the step before it — language models. No step is wasted.

So the trajectory of AI's intelligence is traceable. This year's step is Agents, but even the Agent step will eventually run its course. It'll solve all the problems it can solve, but it still won't be able to replace your employee. It will have reached the limit of its capability — just like CoT. CoT, at its limit, surpassed the very best humans at math Olympiad problems and programming. But it still plateaued there. That particular technique didn't get us to AGI.

So you see, the arc of AI intelligence is traceable. After Agents, the next problem we think needs solving is continuous learning — how to enable models to learn continuously, rather than requiring intense, discrete training sessions. They should be able to sustain learning over an extended period, much like a human.

This problem is of the same nature as task completion and related challenges — they're connected, targeting the same underlying issue.

LIANG WENFENG · CONTINUED

FOUNDER

From where we stand now in the Agent phase, what's visible as the next bottleneck is continuous learning. That's the next problem to solve — how to achieve continuous learning. It's visible, it's relatively clear. It's the obstacle looming ahead that you have to overcome, and there's definitely a way. It just takes time.

After continuous learning, we may arrive at a Singularity. The Singularity being: once the model can learn continuously, it will be capable of doing everything humans can do. It will be able to develop its own next version, conduct its own research, and build the next iteration of artificial intelligence on its own. So it reaches a point — a Singularity — where self-iteration becomes possible.

But this Singularity isn't really a single point. It's a gradual process. The process could be a relatively long, continuous evolution — it's not a sudden leap. We've just gotten used to calling it a “Singularity.” Long ago, futurists predicted there would be a Singularity here, but in reality it's not a singularity — it's a continuous trajectory.

After this step is completed, I believe comes embodied intelligence.

This is our conjecture. Our projected timeline is: first solve learning-to-learn, then reach the intelligence Singularity — the self-iterating Singularity — and only then, embodied intelligence. Once we get to embodied intelligence, it steps into the physical world — it can do your housework, take care of you in old age.

We see this as a relatively ideal roadmap. Different people may see it differently — there's no right or wrong here. It's just that we think this roadmap is the most effortless. In this roadmap, at each step, the amount of genuinely new work is minimal. On this roadmap, we don't need to work overtime.

But if the roadmap were reversed — if embodied intelligence had to come first — that would be grueling, back-breaking work. We don't want that kind of roadmap. We prefer to keep things light. If we first solve continuous learning, then the self-iterating Singularity, and then embodied intelligence, the journey becomes relaxed. Because later on, you can use earlier

technologies to help develop later ones. After the Singularity, when it comes to embodied intelligence, we won't even need to do the work ourselves — the models will just handle it.

So this answers the question of what our long-term goal is. As I've said, this is our long-term goal — what we call AGI.

LIANG WENFENG · CONTINUED

FOUNDER

Now let's come back to reality. Last year's dominant reality was that everyone had to build a chatbot and fight for consumer traffic. This year's reality is that everyone is competing for enterprise revenue, trying to get a piece of that pie — because if you don't, you're not even on the table, right?

But we don't see it as a critical matter. Internally, what we truly care about is exactly what I just described — the AGI roadmap and how to achieve the next technological breakthrough.

There's a strange thing, though: the things you want most desperately are the very things you can't get. The things you don't care about so much — those actually come fairly easily.

Herein lies a strategic advantage: we have AGI in our hearts, and we're building toward AGI. So when we then work on applications — consumer or enterprise — we don't need to invest all that much mental energy. It takes very little effort.

I think there's a “higher-dimensional strike” dynamic at play — when you stand at a higher position in the technology stack and work on something one level down, that advantage is real. On last year's consumer side, we certainly saw it play out. We didn't spend much effort on the consumer side. At one point, we didn't even want to maintain those users — but the users just wouldn't leave. They literally couldn't be driven away. So they're all still here.

And this year's enterprise revenue — the growth trajectory looks fairly optimistic right now. I think the numbers, compared to peers, should be pretty good, I'd estimate. But we haven't put significant effort into it. We basically haven't done anything special — it just happened organically.

As the intelligence ceiling rises on the path forward, the AGI step is one we must take — it's a mandatory rung on the ladder. On the road to AGI, I have to pass through this phase. So I offer all these technologies to everyone via API. I haven't done anything extra.

We're still doing AI. This is a byproduct. I just need a handful of people to maintain the API — no customer support even, no sales, nothing. Users just show up on their own.

You could say, the users we think of as consumer-side or enterprise-side — they're all byproducts of our AGI journey. They're interim outputs. There's no conflict with building AGI. I'm not building the consumer business for the sake of a consumer business, nor the enterprise business for its own sake. We're building AGI to build AGI, and these happen to be outputs we can monetize along the way.

LIANG WENFENG · CONTINUED

FOUNDER

This is different from other companies. Other companies build models specifically to serve consumer users or enterprise customers. But for us, that was never the original intent. Our original intent is still the pursuit of AGI. I think this, to some extent, gives us a higher-dimensional advantage.

AGI is a bigger vision. A bigger vision can rally more outstanding people — it has stronger cohesive power. So I have an organizational advantage. And I leverage that advantage — it's a kind of dimensional advantage.

If you're a commercial company and your vision is serving consumer users well, that's a different story. You'd have other advantages — in product, in user experience, in traffic. But you wouldn't have the technological edge.

The favorable dynamic right now is that model technology matters most. Get the model right, and everything else falls into place. This also explains our developmental trajectory. We genuinely chose AGI. I was never thinking about accumulating lots of users. When we suddenly went viral last Spring Festival — that wasn't in our script at all. We never imagined it. We just wanted to get the technology right.

But at that point, I discovered that our organization actually had an extra advantage over organizations built purely around commercial product goals

— in talent and in organizational structure. That was surprising. At a time when everyone on the consumer side was fighting tooth and nail, the prize ended up going to someone who wasn't even fighting.

This really does validate the logic I laid out earlier — there truly is an advantage in talent and organization. The talent advantage doesn't mean our people are smarter. It's about how these talented people are organized, how they're motivated, and how they collaborate. That's where the advantage lies.

You can gather smart people together, but that doesn't mean they'll naturally cooperate, or be full of passion chasing a goal and getting things done. You need a vision. The experience we've had so far tells me: the AGI vision is very powerful.

OK, that addresses that.

Next question: what are our core interests? I said earlier that we have to be very restrained in many areas. So what are our core interests?

LIANG WENFENG · CONTINUED

FOUNDER

In truth, we only have one core interest: maintaining the stability of our team. That is our biggest core interest — you could even say our only one. As long as I can maintain team stability, I will definitely succeed. Definitely achieve AGI. It's that simple.

As long as people don't leave and we can keep working on it, I'm confident there's basically no risk. It's just a matter of sooner or later. If we hit setbacks and everyone stays, then we just keep going. Money is certainly not the problem. Resources are not the problem. Everything else is relatively easy to obtain.

For us, there is only one core interest — one thing we absolutely cannot compromise on: we must maintain team stability. This is also one of our biggest challenges, or arguably, our biggest risk.

Of course, this risk was substantially mitigated through our recent funding round. The option packages people received are fairly substantial — the amounts are quite large. From a team stability standpoint, as long as the

most critical employees and the longest-serving ones remain stable, the rest are unlikely to leave. Even if their options are smaller or their compensation is lower, they won't go. They're not purely in it for the money. Everyone wants to work in an environment where AGI can actually be achieved. So there's a genuine pull for talent.

Historically, our talent turnover has been relatively low. Compared to peers, our talent churn has always been relatively low. But this remains our biggest challenge — our only challenge, you could say. Everything else is just a matter of time. At worst, other things might set us back half a year or a year, but they won't stop us from getting there.

We certainly won't run out of money. We certainly won't run out of resources. Those aren't concerns.

So much of what we do is designed to preserve team stability. Other than that, I think everything else is something we could do without — something we can exercise restraint on.

We've always been very restrained, unwilling to become adversaries with any internet giant or startup. I want to empower them, or I hope I can assist them in doing what they do. I hope I can help everyone achieve their goals. That's also part of our commercial ethos, as I described earlier.

The premise is that everyone just... Under this premise, we are very willing to assist and help anyone — even our competitors, including Alibaba, Zhipu, and Moonshot AI — to do better. Because we don't lose anything. We're open source to begin with, and open source means, to the extent possible, clearly explaining how things work so you can reproduce our results. If you can't reproduce it, say so, and I'll tell you how to do it. That's inherently part of open source. We don't treat competitors any differently.

LIANG WENFENG · CONTINUED

FOUNDER

Of course, if you're a partner, we'll do even more. But on the big-picture interests, there's no conflict.

In our external engagements, our stance is: we focus only on the AGI mainline. As I mentioned — GPT, CoT, Agents, and so on — we stick to the mainline.

The AI field is vast. There are many things we believe are not on this mainline. For example, 3D, video generation — I don't think they bear much relation to the core intelligence mainline, so we won't pursue them. Then there's world models — I think at the current stage, they're also not closely tied to the intelligence ceiling, so we won't pursue those either.

But if others work on them, we're very happy to help. Whether we have the bandwidth is one thing, but there's no conflict of interest.

We also hope these AI technologies can be applied in various production environments and improve society's productivity — help all industries and sectors increase their efficiency. We're very motivated to do this. Whether we have the time, the headcount, or whether our people are personally interested — that's a separate matter. But in terms of interests, there's no conflict. We want to achieve this outcome.

And we believe this doesn't conflict with business at all. The benefits we're entitled to — we haven't lost any of them. I think the attitude we've maintained so far hasn't caused us to lose out on anything. Our open-sourcing, our goodwill, our willingness to help others — none of it has cost us anything.

Take last year's consumer users, for instance. We still have quite a lot of consumer users now, and they're fairly sticky. This year's enterprise business — I think it's also looking quite positive. Our goodwill hasn't hurt our commercial interests at all. None. If anything, it may have helped.

This seems counterintuitive, but it's actually how things work. Or let's flip it — if we went against this principle, would we be able to capture more? No.

There's a question: how to understand the claim that world models aren't related to raising the AI intelligence ceiling? What I'm saying is, at this current stage — that's our assessment. We have our own AI roadmap. It's not the only roadmap.

LIANG WENFENG · CONTINUED

FOUNDER

From our understanding and our assessment, what matters most right now is doing AI training well. Doing AI training well doesn't require world models — doesn't even require multimodality. If you narrow the scope of AI training a bit and skip multimodality, there are only certain tasks you can't do, but it doesn't affect the validity of the algorithm. Multimodality will have to be done eventually.

What matters now is training. Then the next step is solving continuous learning. Then the next step after that is the model asking its own questions. But this roadmap doesn't include world models or video generation.

When video generation first emerged, it was huge. It felt like something you had to do — if you weren't doing it, you weren't a real AI company. I found that strange. If you just think it through carefully, it has very little to do with the intelligence roadmap. And as it turned out, after Sora first came out and everyone — big companies and small — rushed to do video generation, the small companies later all cut those efforts. It's not related to the intelligence ceiling.

Commercially, though, it's a good business. Commercially, it's a good business. But it has nothing to do with intelligence. We won't do something just because it's a good business — we'll only do it if it's on the intelligence roadmap.

Video generation is a relatively clear-cut example, so I use it to illustrate. As for world models — the definition is less precise, because many things could be labeled a world model. By our assessment, world models and intelligence are not yet the most important issue at this stage. The most important things remain AI training, and after AI training, how to solve continuous learning.

That's our company's assessment. Of course, every company's assessment is different.

As I was saying, the most important issue for our company is personnel stability. From another angle: what are we lacking? What's the gap between us

and the US? The gap really only comes down to one thing: resources. We don't have that many GPUs. Our GPU count is still relatively small.

We currently have about 20,000 H-equivalent units of compute. Most of these arrived recently — in the last month or two. Many machines probably haven't even been racked yet. Our total compute was quite low last year. This year, we're expanding very aggressively.

We're at about 20,000 H-equivalent now. In the coming months, we'll have large batches of additional machines coming in. Almost all of it is NVIDIA.

LIANG WENFENG · CONTINUED

FOUNDER

How many GPUs do we need? Right now, the more the better. Within what we can afford, the more GPUs the better — no question. So our current strategy is: at reasonable prices, buy as many GPUs as we possibly can. Once this round of funding is used up, I'll buy as many GPUs as I can. The pace of spending isn't planned — as long as the price is reasonable, I'll buy as much as there is to buy.

If we burn through all the money in six months, I'd consider that a good thing. If I could spend it all in six months, that would be a blessing — that would be ideal. The reality is that spending this much is incredibly difficult. You can't buy that many GPUs — it's very hard to procure them, and prices are very high. You can't just pay exorbitant prices; you still have to ensure the price is reasonable.

If I could spend all the money in six months, that would probably be the ideal scenario. Because turning cash into NVIDIA GPUs is definitely better than leaving it in the bank. In the bank, right now I'm probably earning around 2%. But buying NVIDIA GPUs gives you a ten-month social cost payback. So you should buy as many as you can.

Even better — after buying GPUs, I have plenty of headroom. Through providing services or whatever other means, I can always generate cash flow. As long as I have cash flow, I survive. I don't need to keep a huge pile of cash on the balance sheet.

So our only worry is not being able to buy enough GPUs. If we could convert all our money into GPUs, we wouldn't hesitate for a second to turn every

last yuan into compute — and we'd be willing to pay a certain premium for it. We are willing to pay a premium to turn cash into GPUs, because it's just such a good deal. Even with the premium factored in, it's still very hard to actually achieve the target.

Objectively speaking, if we could spend 20 billion RMB this year, that would be a superhuman performance by our procurement department.

The gap between us and the US is mainly in resources. The gap in talent isn't very large. There's almost no talent gap, because it's the same pool of people — Chinese people. When Chinese people go abroad, some stay domestically, some stay overseas, some go abroad. It's not the case that the smart ones all go abroad — that's not true. It's actually fairly random. Among the very smartest, maybe slightly more than half go abroad, and slightly fewer than half stay at home. China doesn't lack talent. And we have a large population base — every year we have so many new people joining the pipeline.

Talent is not the bottleneck. Resources are the biggest bottleneck. Resources first and foremost affect talent development — because with less compute, we have fewer opportunities to run experiments, so our talent pool overall lags behind the US. The talent gap, at its root, is also caused by the compute gap.

At the scale of today's largest models, we simply can't afford to train them. Even if we spent all 50 billion, we still couldn't afford it. Even if we could muster the resources to build it, we couldn't afford to run it. The biggest models today have roughly 800B active parameters. In China, we're still operating at the tens-of-billions scale. The largest domestic models probably have active parameters in the tens of billions — that's an order of magnitude difference.

LIANG WENFENG · CONTINUED

FOUNDER

If I wanted to train a model on the same scale as the largest US ones, I'd probably need 50,000 GB300s, or 200,000 Huawei 950s. And that's just for training — not even accounting for research.

So the biggest gap between us and the US is resources. Our current resources — and our resources within this year and the coming months, including the large batch arriving soon — are only sufficient to run more experiments at the 10B active parameter scale. At the tens-of-billions active parameter scale, there are still many experiments to run, many things to figure out.

We're still quite far from being able to train an 800B model. There's still a lot of ground to cover. We just don't have that many GPUs.

So I believe our difference from the US is essentially a resource difference. We might argue that everything we observe — talent gaps, model capability gaps, application gaps — can all be attributed to the compute resource gap. On one hand, domestic GPUs are simply hard to acquire. On the other hand, our capital investment is much lower than the US. We've invested far less at the capital level. Talent salaries account for a very small share in the grand scheme. You see them offering annual salaries of a hundred million dollars — but when you do the math, talent salaries still account for very little. The lion's share is still compute.

This problem is basically unsolvable right now, because Huawei's production capacity is also limited. To train an 800B model, I'd need 200,000 of Huawei's latest cards — and that's just training, not research.

So we're not even considering competing with the US at such a large scale. Our current focus is winning at the scale we can afford to train and run — the tens-of-billions active parameter scale. Get that right first. Then, in the next phase, when we have more resources, scale up to 150B, 156B, or 250B active parameters.

So between us and the US, there's a gap that currently looks hard to close. If you absolutely had to train a model that large, you technically could — but you couldn't do adequate research beforehand. You can't do the thorough research before training.

Another question people care a lot about: in the endgame of the foundation model competition, where will the differentiation ultimately lie? What will the gaps amount to?

I think in the end, the differences won't be that big. The ultimate differentiation will come down to three dimensions: cost, time, and user experience. Beyond those, there probably won't be much differentiation.

LIANG WENFENG · CONTINUED

FOUNDER

Cost is straightforward: for the same service, delivered at the same quality, what cost can you deliver it at? Take the same service — say, BYD batteries. At equivalent technology levels, can others deliver at the same price? I think that's a genuinely hard thing. Not easy to pull off. It's definitely a moat. So cost is certainly a differentiator — I'd rank cost as the number one differentiator.

Second is time — when can you achieve it? Being a few months early versus a few months late makes a difference.

Third might be experience. User experience still varies somewhat. There's still some user stickiness and user-side moats, but that's probably not fundamental. The fundamentals are, first, cost; second, time. Whether you get there first or later — whether you reach the same outcome faster or slower.

As we further enrich our long-term commercialization pathways and product lines, how do we approach pricing? We believe what's most worth doing and worth investing energy in right now is still AGI. Push AGI forward — keep raising the floor and ceiling of intelligence. At this stage, that's a better bet than building more product lines or thinking about more commercialization paths. It's a higher-return path.

I think for some time to come — and indeed for any period in the past — this has probably always been true. Let me put it this way: if we had spent a lot of time thinking about diversifying our products — would it have helped? What commercialization approach would we have brainstormed six months ago? It would definitely have been: I need to do advertising, then I need to do e-commerce, I need to embed e-commerce into the product, and then tie it all together with local services. Definitely useless. Because things change too fast.

If you're at the front end — or at our current stage — and you spend a lot of time on commercialization planning and product lines, those product lines will have very short lifecycles. I don't think we're at that stage yet.

That's our assessment. Or at least, all the experience so far supports this assessment. Over the past three years, at any point, if you came to me talking about commercialization paths and product lines, it would have been a waste of time — because you can't foresee the future. What you can foresee is very little.

Looking at the time — does anyone need a break? Want to grab something to eat? Or shall we continue? No objections — then I'll keep going.

On doing world-leading AGI research and commercializing at the right time. I think we've always been doing commercialization. I believe we've always been commercializing — just not with commercialization as the objective. We have AGI as the objective, but we've been commercializing all along. That's why we have consumer users and enterprise revenue.

LIANG WENFENG · CONTINUED

FOUNDER

Historical experience says this strategy works. And I think the point at which we pivot fully toward commercialization is still very far away. The biggest lever is still technological progress — building the next generation of technology and solving more of the current problems. The ROI on this, in my personal view — at any point in the foreseeable future — focusing on product is simply too early.

So this is also part of our restraint. I hope these commercial opportunities can be seized by others. I hope that how to use AI commercially is something the whole of society, all our partners, can work on together — everyone sharing in the benefits. I don't want to monopolize it. That's impossible. And we don't have the bandwidth — our organization doesn't have enough people to do all of that.

Regarding our partners — we were very deliberate in selecting investors for this round. As for the specifics of how the cooperation is structured, I think first and foremost, our interests are fairly aligned. They're the ones most

aligned with us, least hostile toward us — or rather, the ones who most want us to succeed.

Not everyone wants us to succeed, because we do harm the interests of quite a few others.

On major strategic decisions, technical and business decisions — the process and decision-making mechanism. Our company is fundamentally built on consensus. It's not that I alone decide everything. I seek consensus. My authority and influence within the company are built on consensus.

For example, if I want to do something, I first look at what the consensus is, what people want to do. I might offer some guidance or nudge in a direction, but that guiding influence is limited — very limited. It still rests on a foundation of consensus. The decision-making mechanism is essentially a consensus-seeking mechanism. It's not that I can push through whatever I want. It has to be consensus before I can drive it forward — and only then will I try to drive it.

My main focus right now is basically all on DeepSeek.

HOST

HOST

Let's take a five-minute break.

INVESTOR

INVESTOR

I'll write this down quickly.

HOST

HOST

Feel free to unmute and give me some feedback.

INVESTOR

INVESTOR

We're good on our end. How about we take five minutes.

CHAPTER 02

Strategy and industry outlook

Model competition, organizational focus, the AGI roadmap and the role of cost.

LIANG WENFENG

FOUNDER

When evaluating how different models ultimately pull apart in performance, it should be assessed holistically. Comparing model performance has to be done at equivalent cost for it to be meaningful. It's like comparing two cars — you compare cars in the same price bracket.

The difference between a well-built model and a poorly-built one shouldn't be confined to one specific aspect. It's going to be across the board.

Is Anthropic surpassing OpenAI a lasting shift? I don't think it's lasting — it's definitely cyclical. OpenAI and Google will, in all likelihood, continue to trade leads — they'll rise in alternation. In reality, Anthropic's edge in Code Agents isn't that decisive. It's not as if they're crushing OpenAI. In our company, roughly half the people, at any given time, feel that OpenAI is better.

Anthropic had a first-mover advantage, but that advantage should fade quickly. It's not something they can hold onto for long. All three of those players are incredibly strong. Among the three, Anthropic is actually the most efficient — the amount they've spent, the money they've burned, is probably the smallest.

When it comes to the division of labor in global AI, the role Chinese companies are most likely to play is still being the largest-volume producer. By common logic, we have the largest production capacity — including for chips — and we have the most electricity. So ultimately, our AI is very likely to be one of the three pillars.

Chinese people will make these products the cheapest while delivering comparable results. After all, for many products today, the gap between

Chinese-made and American-made is not that big. The future of AI might follow the same pattern — but Chinese AI will probably be cheaper. This cheapness may be systemic, just like how in other industries, Chinese services are cheaper. It's probably the same dynamic.

When I think about what to do, my habitual framing is: at this moment, what course of action yields the greatest return? If I felt that building products right now yields the greatest return, I'd go build products. If I felt that achieving AGI first yields the greatest return, then I'd focus on AGI. Obviously, I don't think building products right now yields the greatest return.

On the question of domestic GPU adaptation — domestic chips actually face a historic window of opportunity right now. Previously, adapting domestic chips had a core difficulty: the ecosystem was poor. After buying the chips, you couldn't really use them — they lacked NVIDIA's ecosystem. That meant NVIDIA's moat was very deep.

But this is changing. NVIDIA's CUDA moat is being rapidly dismantled. There are probably three driving forces behind this rapid erosion.

LIANG WENFENG · CONTINUED

FOUNDER

First, we now have AI. And with AI, building out an ecosystem has become far easier than before — because AI can write code. I can use AI to construct the ecosystem, essentially replicating NVIDIA's entire ecosystem. That's the first factor: AI.

Second, there are new technologies. For example, we've developed a technology called TileLang — it's a high-level language. Using this high-level language to write CUDA kernels, you can very quickly rewrite NVIDIA's entire ecosystem from scratch. Combined with AI, this looks like a path with no real obstacles. It hasn't been completed yet — it's not done — but the technical route appears to be unimpeded.

There's another point. CUDA — NVIDIA evolved it out of gaming GPUs. So in many ways, the design of gaming GPUs and the design philosophy run in a continuous lineage. CUDA maintains backward compatibility with gaming GPUs. In the past, AI compute was a small niche — smaller than the gaming

GPU market — so that made sense. But now the compute card market has surpassed the gaming GPU market. There's no longer a reason for these two to remain coupled.

The trend now is toward decoupling. Dedicated AI chips — whether from Huawei or even NVIDIA itself — in the future will be purpose-built silicon, not extensions of what came before. Against this backdrop, the relevance of NVIDIA's legacy ecosystem is dramatically reduced. Because for a dedicated chip, there's very little relationship with CUDA — they're not bound together. Or rather, the chip is designed from the ground up with the ecosystem in mind.

It's a bit complex. But anyway, domestic AI chip substitution is facing a historic opportunity right now. We believe that within a year, we'll see one thing validated: the ecosystem for domestic chips is completely fine. Previously, people thought it was problematic — that they couldn't be used, weren't practical. But I believe within a year, we can reverse that perception — or rather, facts will reverse it.

The hardware and ecosystem of domestic AI chips are both fine. The only problem is insufficient production capacity. On the adaptation front for domestic chips, there are no barriers. NVIDIA can't stop it.

In a normal commercial environment where you can buy NVIDIA GPUs, domestic substitution would be relatively difficult. But when NVIDIA GPUs are unavailable, everyone is forced to pivot to domestic chips. Against that backdrop, domestic chip adaptation faces no obstacles at all.

Building a domestic chip ecosystem on par with NVIDIA — or even better — I see no obstacles. It just takes time.

Our main collaboration right now is with Huawei. They're doing their own adaptation, but we get involved in the ecosystem work — we're deeply involved on the Huawei side. Huawei's problem is still capacity.

LIANG WENFENG · CONTINUED

FOUNDER

Huawei has offered us roughly 16,000 cards of capacity. The big internet companies might get a hundred thousand-plus. We get a bit over ten thousand — I think this ratio is... but this is probably all the capacity Huawei

has. So we can't rely on Huawei to train the next larger model — or to train models with several hundred billion active parameters — certainly not this year. But next year, or the year after, there might be a real chance.

On Huawei chip adaptation, our main work is getting the high-level language compiler right — getting TileLang right. Once TileLang is done, the problem likely solves itself. This is somewhat complicated to explain, but we're working on it. Once it's done, it'll be much clearer to describe.

You can think of it this way: when V3 was trained, it still used NVIDIA hardware — but it no longer used NVIDIA's ecosystem. V3 used NVIDIA GPUs without using NVIDIA's ecosystem. Instead, we first wrote a high-level compiler called TileLang, then built everything else on top of the TileLang ecosystem — almost entirely independent of NVIDIA's ecosystem.

As long as I replicate this entire stack — do the same process over again on Huawei hardware — it'll be done. I think this might be a historic mission: it could completely overturn the previous perception that the domestic chip ecosystem is poor.

A lot of the stars are aligned now — timing, circumstances — all that's missing is time. I think within a year, many people will be on board, or at least will understand that this problem is essentially solved. What remains is the capacity issue.

I'm relatively optimistic about domestic compute. On this point, I'd say NVIDIA is digging its own grave.

Huawei's SuperNode — the Huawei 950 SuperNode — can fully replace NVIDIA's GB200 and GB300 in both performance and price. The price will definitely be higher, but the premium is limited. Fifty percent more, a hundred percent more — even a hundred percent more is fine. Two hundred percent more is also fine. If it's a hundred percent more, I'd already consider it price-parity in practical terms.

In terms of tasks, it's also a drop-in replacement. Everything the GB300 can do, the Huawei SuperNode can do — latency and everything, the same. The only cost is: four Huawei cards equal one NVIDIA card, and they're two years behind.

Four-to-one is understandable. “Two years behind” means four Huawei 950s match one GB300. The time lag is two years. The Huawei 950 SuperNode is shipping in Q3 or Q4 this year. NVIDIA's GB200 shipped in Q3 two years ago — so there's a two-year gap. By Q3 this year, NVIDIA may already have its next generation out.

LIANG WENFENG · CONTINUED

FOUNDER

So the chip gap between us and the US — I believe that in the future, there won't be an ecosystem gap anymore. But on the chip level, it's four times plus two years behind.

Another question: will we vertically integrate upstream? I hope we don't have to.

There's a question: are we planning to vertically integrate into upstream applications? I hope not. I hope others take on that work. I don't want to eat up everything. I want to eat just one slice — just the slice I'm best at, the slice we consider the most core. And the slice directly touching users — I think for many of our industry partners, they care more about that than... It should be their domain, not something I appropriate.

Will we build our own large-scale clusters in the future? I think building our own large clusters is a must — we've been doing it all along. All our clusters are self-built. But whether we'll get into chip development — that depends on how large the payoff would be.

Tesla... power generation. Here's an analogy: if you operate a power plant, do you necessarily need to build your own generators? Power generation equipment can be made by others. If the price is reasonable, why would you build it yourself? So I hope we don't have to make chips. I hope we can buy chips at reasonable prices, so I don't have to go into chipmaking.

I think this is very likely — that the massive profits in NVIDIA's chips right now may not be... Although... we hope to do just one piece.

I think AI is big enough that I don't need to... I'll just do one slice. If I stay focused — and I believe the business upside here is already enormous — if the AI era is going to spawn multiple trillion-dollar companies, I think we'll be one of them. We've been doing just one small piece, and being one of

those companies is already not in question. I don't need to... I'm not ambitiously trying to be everywhere.

I think the most likely outcome is that it's hard to shift focus. And if you really do want to branch out, you'll still have more ideas — this will... This is our attitude.

For example, in both enterprise and consumer businesses — at least from what we've seen so far — the companies that genuinely wanted to close the consumer loop ended up not doing it as well as we did. The companies that genuinely wanted to close the enterprise loop may also not do it as well as we did. The more you want...

And on the enterprise side, let me say a few more words. The ceiling of the enterprise business should still be demand. Under the current generation of AGI and AI technology, enterprise demand is probably finite. It will grow rapidly, but it's not an infinite thing. Ultimately, it's constrained by demand, not by compute.

LIANG WENFENG · CONTINUED

FOUNDER

Under current technological conditions, how large the revenue can get ultimately depends on demand. Think about it: I can recover cost in ten months — so if there were uses for it, I'd buy more. But there just aren't that many use cases.

Right — demand should keep growing. If technology continues to make breakthroughs, demand will keep growing.

On multimodal — we've been working on it all along. For products, it's very important. For consumer-facing products, it's very important. But for the intelligence ceiling, it's a component — it's not the mainline itself. However, as a component, multimodality is something we'll definitely do — and we are doing it.

We should be releasing related models — our V4 and subsequent versions will support native multimodality. But when it comes to multimodality relative to intelligence, we treat it as a component, not as intelligence itself. Its

role on the mainline is similar to search. Search is also a component; multimodality, in our understanding, is also a component.

On scaling — we believe in scaling. Larger scale definitely yields better results and unlocks more capabilities. What's stopping us from scaling is really just compute. It's not that we don't want to scale — we just don't have enough compute to do it.

We haven't reached the ceiling yet. The scale at which we train our models isn't because we think that scale is sufficient — it's because that's exactly how much compute we happen to have. I work backward from my resources: what scale of model can I afford to train? That's how it's calculated. It's not that this model size is “enough.”

As things stand, the returns from scaling are still very clear. We haven't yet had the chance to hit the scaling wall. That's still far away from us. When Silicon Valley talks about scaling hitting a wall, that's from Silicon Valley's perspective. For Chinese players, we're still very far from that — we simply haven't scaled to that degree. This scaling includes data scaling, model-size scaling, and training cost scaling.

We're still fairly far from exploring the upper bound of scaling — because we don't have that much compute. But we will spare no effort in pushing that scaling frontier. After this funding round, with more compute, we'll probably train larger models. The goal is to buy within the shortest time possible. If we could burn through all the money in six months, that would be optimal — but in reality, it can't be done.

On the core capability of the next-generation model — I believe it must possess continuous learning ability before it can be called a next-generation model. Before that, what we can do is: drive cost down, make performance better, make inference faster. But for a major breakthrough, it has to have continuous learning.

LIANG WENFENG · CONTINUED

FOUNDER

Another question: why does it seem like we care especially about model compute efficiency? Because I've indeed observed that not everyone cares very much about model efficiency. For a commercial company, there's no

strong incentive to pursue model efficiency — a bit more efficiency doesn't change much. So there isn't that much motivation to push efficiency to an extreme.

So several startup companies — you don't hear them say that low cost is what they're pursuing, because that doesn't align with their interests. Low cost — then what do you make money on? Low cost means you can't charge much.

Conversely, this is part of our vision. Or rather, our vision — our people care about this cost, because our people are ordinary people and they know that using this stuff costs money. They can empathize: other people have to pay to use it, so if it's cheaper, the barrier to adoption is lower. So many of our team members genuinely want us to drive cost even lower.

But from a purely commercial standpoint, you wouldn't think this way. Commercially... from a cost or commercial perspective, this isn't the top priority. Whether for a startup or a big company, the cost of providing this service isn't high at all.

But I think we want it to be relatively light — I want it to be affordable. Especially in China's context of compute scarcity, it should be affordable and usable on domestic GPUs.

I think low cost is, first and foremost, a result. Our models have indeed been consistently heading in a lower-cost direction architecturally, and this is connected to our vision. We still have plenty of algorithmic approaches that can drive cost down further.

Another reason to push cost down: the lower the cost, the bigger the model I can train, the bigger the model I can afford. With the same amount of compute — with compute constrained — if my compute efficiency is higher, I can afford a larger model. For a big company, they might not think this way. For a big company, resources can simply be added — solved by throwing more resources at it. But we prioritize cost efficiency.

On the value of data... data covers a broad scope. Data is essentially half the model.

Why do I think that if I wanted to — say if AI could account for 20% of GDP, and then I wanted to take 5% of that — that would absolutely not work? Because I'd inevitably be defeated by someone else. If there's another player who says, "I only want 1%," then I'll surely be beaten by them.

LIANG WENFENG · CONTINUED

FOUNDER

If my goal is to capture 5% of AI, or 5% of all of humanity's GDP — theoretically, you can make the math work. Look at OpenAI: they can make that math work on paper. Theoretically, no problem. But the issue is: they'll be defeated by someone who's willing to take only 1%. Because the other player says: I can deliver just as well, but I only want 1% of global GDP — and that'll beat them. At that point, yet another player might emerge and say: I only want 0.1% — and that'll beat the previous one.

At the macro level, regardless of where that percentage is taken from, there's no difference among them — it's all the same. Those who take more get beaten by those who take less. You don't even have to actually take more — if your vision is to take more, you'll be beaten by someone whose vision is to take less. Even before anyone has actually taken any money — just the vision alone. If your vision is to take more, you've already lost. You'll face greater difficulty.

That's just how the world works. OpenAI, from the start, felt they could genuinely monopolize this world. But in reality, they'll face challenger after challenger. They'll face challenges, so they won't have it easy. The US will face challenges, and in the future, it'll also face challenges from China — because Chinese people are willing to take less to provide the same service.

Within China, there will also be people willing to take a bit less. But eventually, there'll be an equilibrium — because if you take too little, the company's business logic no longer holds and you can't survive. So take too little, you die. Take too much, and someone taking less will beat you.

So for us, it's not about maximizing profit — or pricing for maximum return. It's about earning only a reasonable return. That's one way to explain it.

I genuinely believe this. I'm not rationalizing after the fact — there's no need to rationalize. I've been doing it this way all along, so of course there are reasons. They might not be the most conventional reasons, but I think this company was never conventional to begin with.

Our management runs on two tracks: one top-down, one bottom-up. Bottom-up means: everyone does what they want to do, on their own initiative. No one manages them. No KPIs. Top-down means “official business” — when we need to do something collectively that requires company-wide coordination. For example, when we need to ship V4, we divide the work, and everyone takes a piece of it. That's top-down. And we call this top-down work “official business.”

Generally, we want “official business” not to take up more than half of each person's time. “Official business” shouldn't exceed 50%. The other half of their time is unscheduled — they do whatever they want.

This is in the research domain: letting people explore on their own, pursuing whatever they think is important, with no preconditions. As long as the company can support it — as long as there's compute available to do it, or they don't need much compute, so they don't even need to coordinate — they can just go ahead.

LIANG WENFENG · CONTINUED

FOUNDER

So this is how we're organized right now. Some people think we're top-down; some think we're bottom-up. I think both are right.

One rule of thumb I have: “official business” should preferably not exceed half the time.

We generally don't work much overtime. There are two reasons. First, research requires a relatively relaxed environment. If you push too hard, you can't do research. Because it's about people having genuine interest and thinking about these problems on their own time. So it has to be in a relatively loose, relaxed environment for exploration to be possible. That's a necessity of research culture.

Second, we're very focused. Being very focused means we have very few things to do. So I just don't have that many things to do — which means no

need for overtime. This is all of a piece with the restraint I mentioned earlier. Because I practice restraint, I just don't do many things. With fewer things to do, each person's workload is lighter.

You can see that many of our products aren't polished — and we haven't gone back to fix them. That's also part of our culture.

OK, I've skimmed through most of the questions. If anyone has more, go ahead.

CHAPTER 03

Continual learning, talent and compute

The first Q&A block: research bottlenecks, infrastructure, domestic hardware and the China-US gap.

HOST

HOST

Investors, please feel free to unmute and chat. A quick reminder: Liang Wenfeng has shared quite a bit of sensitive information. Please do not leak any numbers or specifics, including GPU counts and the like. Also, absolutely no screen recording or external sharing. Thank you all very much. If you have questions, please unmute and speak freely.

INVESTOR

INVESTOR

Mr. Liang, could you share more about the timeline for when continuous learning might lead to a breakthrough? And what architectural or algorithmic innovations would be needed to achieve it, along with other key elements?

LIANG WENFENG

FOUNDER

It needs fewer people and more research. Right now, the entire world is researching this problem. Or to put it another way: from an investor's perspective, what investors see most right now is Agents. But for us researchers, what we see more right now is learning — and how to solve the learning problem.

Actually, learning may not be a single technique — it's a problem. How to solve this problem could involve many different techniques — not just one technique, not a single thing. It will be many things. AGI is composed of many elements — you need models, and then you need a lot of other pieces. It's actually an engineering and algorithmic challenge.

This question is quite specialized, but there are many approaches and a lot of ongoing research.

INVESTOR

INVESTOR

Mr. Liang, thank you. We're very grateful for this opportunity today.

First, I want to respond to and express appreciation for what you said at the very beginning — it moved me deeply and gave us tremendous inspiration. You mentioned that this team, carrying the deepest goodwill, hopes to contribute meaningfully to this industry, to society, and to the advancement of human intelligence — to make a small difference. And you do this with a sense of mission and vision.

I think this closely aligns with the corporate culture we serve — the principle of “cultivating oneself to benefit others.” I've come to deeply understand why you lead the team in embracing open source. I picture it vividly: like building a banyan-tree ecosystem — a paradise for all birds — benefiting all things without competing with any of them. In that way, it is welcomed by all; everything coexists and co-thrives, and ultimately, it becomes omnipresent.

So through this investment, we express the same recognition, support, and respect for this mission and vision. At the same time, we hope to contribute in the areas we're strong at as this industry develops.

And on this, I'd like to continue the conversation and seek your thoughts. For example, in the future co-building of the ecosystem — now that things are open-sourced, how many partners, talents, and teams in the industry are capable of reasonably reproducing your open-source models and results? As the ecosystem evolves further, in which areas do you feel more high-quality talent is needed to connect with your models and reproduce them?

Or is it that GPU compute is currently somewhat scarce? Will the future be some kind of model-matrix approach? For instance, your foundation models keep getting better and better, while partners and teams across industries build vertical industry models or application models within that

matrix. How is that developing at the moment? Step by step, over the next two or three years, what kind of ecosystem do you envision taking shape?

That's my first question for you.

Second: you also shared quite a bit with the group about observations on AI hardware. Global AI giants may individually invest at the scale of hundreds of billions of dollars. China currently appears to have some shortcomings in hardware compute. How long do you think it will take to resolve this and support our AI development so that the compute and hardware shortfall doesn't hold AGI back?

Do you feel this is a mission that Chinese people must and will accomplish — that we can definitely do it, and it's merely a matter of time and capital investment? But at the same time, it could be a two-sided dynamic.

INVESTOR · CONTINUED

INVESTOR

On one hand, as models improve and intelligence advances, the compute and hardware requirements for a single task or a single intelligent agent may gradually decline per unit — no longer requiring such massive compute, because smarter models enable clever computation. I wonder if my understanding is correct.

On the other hand, hardware advances will make compute power more efficient. Might these be two paths converging from opposite directions?

If we're at this point in time, using the current 960 or H200 to build hundreds of billions of dollars' worth of compute, you mentioned amortizing over three years. What do you think its actual lifecycle is? Will technology iterate in four or five years?

Or more bluntly: if a compute center now builds a 10,000-card cluster with today's GPUs, in three years might it be relatively obsolete — second-tier compute? Could we see a situation where right now it's under construction and insufficient, but three years later it becomes suboptimal compute sitting idle? I wonder if such a phenomenon could occur.

Two questions for you. Thank you. Thank you.

LIANG WENFENG

FOUNDER

First question is about the ecosystem. Our current sense is that every company is probably facing a talent shortage. But I think this talent short-fall will be temporary. At the early stages of every industry, there's never enough talent. Think back to building websites — when websites first started, there were very few people who could build them; talent was scarce. Later, when the internet moved to server-side development, talent was also very scarce. But this kind of talent shortage gets resolved very quickly — within two or three years — because large numbers of people get trained up.

The shortage of AI talent is also temporary, and we're already seeing it significantly alleviated. Because there really is no shortage of AI-capable people. Every company quickly trains its own people — training people is fast. So in the AI industry as a whole — whether ecosystem, model companies, or whatever — talent is not in short supply. Talent shortage is definitely a short-term phenomenon. History has never seen a persistent, long-term shortage of any particular type of person.

I still remember, over a decade ago, people said there was a severe pilot shortage — pilots have a long training cycle, but that was also resolved quickly. So no one needs to worry about talent shortages.

And domestically, there are still too many companies building foundation models — too many. The US may have about three. In China, way too many are building base models. Eventually, you definitely don't need that many people working on foundation models — it will converge. Resources are also too dispersed, which, to some degree, is wasteful. Every single player is doing the same thing, while in the US only three are doing it and resources are concentrated on just those three. In China, resources are split very thin — each player gets even fewer resources.

I think this will definitely converge — it has to. But it takes a process; eventually it will converge. You don't need this many players. Right now, everyone probably thinks the profit margins on this are very high, so they have to do it themselves. But once they discover that this may not be that high-margin, they may stop. Lately, it's certainly not that high-margin — I

don't believe it's that high-margin, because that wouldn't conform to objective laws.

This means we're at a stage where: if there's an extremely high profit margin, that definitely doesn't align with objective laws. It should be a reasonable profit.

So this is the current state of the industry — I'm sure it will converge. The segment where everyone is building large models — forget about any single party monopolizing and saying “I want to take all the profits.” That definitely won't work. If every player only takes a reasonable profit, then you really don't need that many people building large models. In China, having three or four players competing is more than enough competition — prices would absolutely reach a level where price wars are already fully fought. Large models — not to say two big companies and two small ones — that's probably already enough.

As for the ecosystem, I don't have many specific thoughts. We hope to support more people, but we don't have that much bandwidth. We have the will, and there's no conflict of interest — but whether we actually go do it is another matter. At least there's no conflict of interest here. We hope for win-win cooperation.

First, I absolutely do not believe that foundation model companies can take the lion's share of the profits. That's impossible, because there are so many foundation model companies, and the gap between them doesn't need to be that big. The gap is only two things: time and cost. So no single player is going to enjoy windfall profits — I don't think windfall profits exist. Those who control costs well earn a bit more; those who control costs poorly earn a bit less. That's about it.

LIANG WENFENG · CONTINUED

FOUNDER

Did that answer your question? Is the first question answered?

INVESTOR

INVESTOR

Can everyone... I believe in the future there will definitely be many... the future actually will... people's data applications and iterative cycles... can

you hear me now? Thank you. Yes, thank you for your answer. I also very deeply understand and respect your ecosystem strategy and positioning in the industry.

For instance, on the data side — currently, publicly available data, I believe model companies already have the channels to acquire it. That approach shouldn't be an issue — just a matter of time and cost. Then later, when we truly reach AGI, one possible vision or ideal state is that models can self-iterate and self-learn — essentially training themselves.

On that front, do you think synthetic data could be usable? Or is real data the highest quality? If it still requires real data, wouldn't that constrain AI's intelligence to the bounds of human history... because it depends on data that humans have genuinely generated in the past? Or is it possible to break through that ceiling — through simulated data, synthetic data, generated data, and other methods — to enable model capabilities to surpass everything that's come from real human data...

LIANG WENFENG

FOUNDER

I think it can surpass. I see two ways it can surpass. For example, in Go, AlphaGo played a move that no human had ever seen. It clearly surpassed humans within a certain scope. But it might also have an upper bound — it might have its own limitations. However, we can't yet see those limitations.

So broadly, we believe that it can surpass — building on the knowledge that humanity already possesses, the knowledge we have already articulated.

INVESTOR

INVESTOR

And will that part rely on real data or synthetic data going forward? Will it work?

LIANG WENFENG

FOUNDER

It can't be just one method — there will be many approaches.

INVESTOR

INVESTOR

OK, thank you. I'll take a bit more of your time — I also wanted to continue the earlier discussion about AI Infra.

LIANG WENFENG

FOUNDER

What was the second question again? A bit...

INVESTOR

INVESTOR

Alright, let me restate it briefly. I wanted to ask about AI Infra. Looking ahead, it seems compute investment is in the hundreds-of-billions-of-dollars range across the board. On that front, we might believe that Chinese people, on the hardware side, are destined to accomplish the mission — that one day there will be high-efficiency compute. But in the process, it's currently still a constraint.

So will the future be two forces converging from opposite directions? On one side, as model capabilities iterate, the compute demands shift from brute-force computation to clever computation — meaning the compute and hardware requirements per unit/model/task will gradually shrink at the margin. On the other side, hardware advances — things like GPU capability improvements — will make iteration faster and single-card efficiency higher.

So what does this dynamic look like in practice? Could it be that building a 10,000-card cluster now using H200 or 960, but after two or three years, it becomes relatively suboptimal compute — relatively obsolete hardware?

LIANG WENFENG

FOUNDER

For NVIDIA GPUs, you can basically use a five-year depreciation schedule. For Huawei GPUs, three years at most. The Huawei 950 is quite usable this year; next year, I think it'll still be OK; beyond that, I think it might genuinely consume too much power.

Huawei's GPU lifecycle is definitely shorter, because it was already two years behind NVIDIA. But I don't think the gap is that huge.

If B200s were available now, I think whatever you could buy would be worth it. For someone like Tencent or Alibaba, if they could buy at a reasonable price, it would definitely be worth it. But this isn't a moment for calculating cost — I believe you can't even buy them.

INVESTOR

INVESTOR

Understood.

LIANG WENFENG

FOUNDER

The fact that we're behind in compute is a reality. This reality gets mitigated through three avenues.

First, we have to accept some degree of model lag — we can only use smaller models. How much smaller is the question of training. We have to accept a certain degree of model lag and smaller model sizes.

But there's an upside to being behind: being behind means you have more time, and you have access to certain techniques. This means you can use clever methods.

INVESTOR

INVESTOR

Thank you.

LIANG WENFENG

FOUNDER

So the gap between us and the US might be about 12 months behind — maybe 12 to 18 months, or 6 to 12 months. To put it simply: we're about two years behind the US, but we accomplish this with only one-twentieth of their compute. That's the narrative — one to two years behind, but using only 1/20th of the compute.

Going forward, we want to rewrite that narrative: using some fraction of their compute, but shrinking the timeline — down to 6 months, 3 months. I think that's a worthy goal. And we may even surpass them in certain areas. But with an order-of-magnitude gap in overall compute, surpassing them

across the board is unrealistic. However, in certain focused, deliberately chosen areas, it's possible that we surpass them somewhere.

INVESTOR

INVESTOR

Understood. Thank you, Mr. Liang. Full of confidence — let's work hard together. I'll leave the time to other partners. Thank you.

Thank you for the sharing just now. I have two quick technical questions. In the technology roadmap you laid out, you mentioned that the core problem to solve at this stage is continuous learning — which is also a hot research topic abroad, called Recursive Improvement.

My first question: right now, what do you see as the biggest technical difficulty? From your perspective, when can this be solved? That's the first.

Second question: you also mentioned solving continuous learning first, and then moving on to intelligence. I'd like to understand — what is the technical rationale behind that sequencing? Does it mean that after solving continuous learning, DeepSeek will also pursue general intelligence afterward? Two questions for you.

LIANG WENFENG

FOUNDER

The technical question is a bit tricky to explain... The difficulty lies in the fact that we haven't yet found a method that really works. Nobody in the world has figured out a good approach yet — everyone is still groping in the dark. So we're still in the exploratory phase — nobody knows who will stumble upon the next method that cracks this problem. We're still exploring.

We have many lines of thought, many ideas that look promising right now, but none of them have been fully worked through yet. That's the first.

Second — internally, we're quite attached to this narrative: when we train our next model, we want it to be able to assist our own development. It should boost DeepSeek's own efficiency. Our models' first priority isn't to be great for users — it's to be great for ourselves. First, useful to us.

Once it's useful to us, developing the next model becomes faster. A lot of people internally think this way: first, it has to be useful to us — first, it's for our own use. And this is the fastest path to AGI. When it works well for us, that likely means it works well for others too — but it has to work for us first.

This narrative is a bit unusual, but a lot of people genuinely do think this way. Not “good for users” — I want it to help us more, so that we can achieve AGI much faster. So the logic of this narrative is: to help us achieve AGI. But first, to help us achieve it. We need this help to achieve AGI. It's now very clear — we genuinely need artificial intelligence to help us achieve AGI. Even though it's not yet working autonomously — it's still working in tandem with humans — it's already very useful.

INVESTOR

INVESTOR

The second question: you mentioned solving continuous learning first, and then moving into general intelligence. Is that your expectation for what comes next? I want to understand the technical rationale behind this — why solve continuous learning first, and then general intelligence? Your understanding here.

LIANG WENFENG

FOUNDER

Because solving continuous learning would dramatically accelerate our R&D progress. If I solve continuous learning first, then general intelligence is no longer a heavy lift. With AI assistance — if AI can learn continuously — its capabilities should be extremely strong.

The current limitations of Agents stem from the fact that they can't learn continuously — they can't effectively engage in continuous learning. If we could wrap up continuous learning first, AI's capabilities would be enormously powerful, and it would greatly enhance the efficiency of our own research.

Do continuous learning first, and general intelligence becomes easy — you can just use it to build general intelligence. That's why I say this is the sequence we most hope to see — it saves us effort, it makes things easy for

us. Otherwise, right now, doing general intelligence manually — it's tiring, grueling, data-intensive, labor-intensive, and the cost-benefit ratio isn't great.

INVESTOR

INVESTOR

Thank you for sharing.

HOST

HOST

A small question — check the chat window for the online questions. How much longer until AGI? Can domestic hardware catch up in that timeframe? It's in the Zoom chat window.

LIANG WENFENG

FOUNDER

OK, I saw that. Huawei 950 — Huawei is currently giving us 16,000 cards. This should be OK to share publicly. It's probably an order of magnitude less than what the big internet companies get. That's all Huawei can give us, and the price isn't cheap either.

The big internet companies' demands are larger — for them, they need it more. For us, we can also buy some non-sanctioned cards. So the purpose of buying Huawei 950 is really to help Huawei get this ecosystem right.

Sixteen thousand Huawei 950s is only equivalent to about four thousand B-series cards. So it's not a very large quantity — the significance isn't huge. It's not enough to train a next-generation model. It's only enough to train our current-generation model — not enough for the next generation. But it allows Huawei to get things in order first. So that's regarding the Huawei 950.

Then: how much longer until AGI? Can domestic hardware catch up in that timeframe? I think in the AI endeavor — in this AI pursuit — domestically, within a year or two, we should be able to reach rough parity with overseas. Or possibly, as early as this year, we could achieve functional substitution for foreign models. In AI, under the current paradigm, it's not an extremely difficult thing. So this year, it should be achievable.

But it's still not AGI. At minimum, I think it needs to be capable of continuous learning.

Can domestic hardware catch up in that timeframe? I think domestically, we may need a few years. First, we have to solve the ecosystem problem, because the ecosystem is a matter of confidence. Solve the ecosystem problem, then solve the capacity problem — I think these can be progressively resolved. I don't really believe that five years from now, we'll still be stuck on capacity. Right now, we're definitely stuck on capacity — this year, next year, the year after, I think we'll still be stuck on capacity. But five years out, I don't think so — I'm relatively optimistic.

Second question: thinking about future organizational structure and headcount planning. First, our previous organizational structure was extremely decentralized — because there was no organizational structure. But as our headcount expands further, we'll certainly need to make some changes.

All I can say now is: a lot of changes will be needed here, but it's hard to articulate them all at once right now. Ultimately, we'll need to split into different departments. Some departments will need to establish fairly rigorous hierarchical structures. Others will probably maintain a relatively loose, flat structure. As headcount grows, we'll make these adjustments.

We probably need to make these adjustments very soon — I'm already working on them. If we don't make these adjustments, a lot of things simply can't move forward. There really are many departments that should have organizational structure.

LIANG WENFENG · CONTINUED

FOUNDER

And then there's a question about CV — which version comes first, right? I think our currently released CV V4 version is still fairly rough, and there's a lot of capability work that still needs time. For me, a comfortable release cadence is roughly one release every two to three months. The last release was around end of April, so the next release should be around end of June — roughly like that.

Barring surprises, each version should be better than the last. At the 50B active parameter scale, I feel the eventual difference between us and the

current wave of open-source models won't be that large. In terms of inference speed and performance, I don't think there'll be huge differences.

But compared to their larger-scale models — the ones they haven't made public — the gap should still be quite significant. That gap is something our active parameter scale probably can't close — it would definitely require a larger model, maybe 150B.

As for 150B — at our current training pace, this year... optimistically, we could start training it by the end of this year, or at least by next April... the gap is still quite large. Yes — that's the gap relative to OCE.

CHAPTER 04

Scaling, self-improvement and embodiment

How gradual self-improvement, language-model scaling and embodied intelligence fit together.

INVESTOR

INVESTOR

Hello, I actually have a question. You often say that the process of achieving AGI is gradual rather than involving a sudden mutation. So can I understand this as a process without a critical point?

LIANG WENFENG

FOUNDER

It has no critical point, but it's nonlinear. We currently believe fairly strongly in this narrative: AI can accelerate AI research. AI can accelerate AI research. Meaning, it's not linear — because you can use AI to accelerate your own research. So further down the line, it may become nonlinear.

INVESTOR

INVESTOR

Understood. So at the moment, my takeaway from your earlier conclusions is: continuing to scale language models is sufficient — sufficient to reach that state.

LIANG WENFENG

FOUNDER

You can only say: for language model scaling, right now I don't see an upper bound. At our current intelligence level — or the intelligence level reached in the US — nobody has seen an upper bound.

INVESTOR

INVESTOR

Understood. Because I'm very curious — you mentioned that for the US, an 800B active parameter model: even if they can train it, they can't really

deploy it for widespread use. So they can only train it — it's very hard to put it in front of users, because it's genuinely a bit expensive.

One thing I was very curious about before: humans only acquired language ability in the last hundred thousand years or so, but evolution before that spanned 3.7 billion years. In training AI, the sequence you described could indeed be reversed. But ultimately, you'll probably still need to enter what might not necessarily be called a world model — but physical models, or embodied intelligence, right? Beyond this upper bound.

LIANG WENFENG

FOUNDER

Right. I think embodied intelligence is ultimately unavoidable — that's the endgame. So for our company, the natural endpoint is probably embodied intelligence.

Because for an ordinary person, their needs aren't about a computer, right? An ordinary person — eating, drinking, entertainment, clothing, shelter, transportation — doesn't need a computer. What they need... so you still need embodied intelligence to solve real human needs. If the goal is to address human labor needs, then embodied intelligence is, I think, unavoidable.

INVESTOR

INVESTOR

Understood. So in stages — suppose we reach a point that's not necessarily a critical point, but the ability to self-evolve, to self-evolve fairly well, or near that point on the ascending curve — I'm very curious: what would be the first application you'd want such an AI to land on? It might be different from the present.

LIANG WENFENG

FOUNDER

What we hope it can do — directly, assuming no embodiment yet — is: what do we define as AGI, or what do we hope AGI can do? It can help me iterate the next model. Help me iterate the next model.

And then once embodiment is there, what we want it to do is also: let it iterate the next version of embodied intelligence — let it build the next generation of robots.

INVESTOR

INVESTOR

I'm still very curious about one thing. From DeepSeek's previous interviews and so on, it seems that for important direction choices and research decisions, taste and intuition are very important — not just simple engineering optimization. If AI can self-evolve in the future, do taste, intuition, and those things still matter — or what would become important?

LIANG WENFENG

FOUNDER

AI doesn't lack taste and intuition right now. What it lacks is the ability to learn continuously. AI's taste and intuition are fine. Ask it to write an essay — its taste and intuition are, I think, perfectly fine.

There were a few more questions earlier — let me check. I saw some questions on the screen, but I can't see them anymore.

INVESTOR

INVESTOR

Mr. Liang, I left a question on the screen. Let me read it out for you again. I wanted to ask: continuous learning — as you mentioned, many researchers also say it's an unsolved research problem. And then coding agents, especially catching up to MILES-level performance, reaching Office-level or MIS-level capability — is a relatively certain target.

For a research scaling goal that's still unsolved versus a relatively certain scaling goal — how do you think research resources should be allocated, especially in terms of researcher talent? How to strike the best balance and achieve the best results?

LIANG WENFENG

FOUNDER

Model performance at the Office level is a relatively certain target. But model performance at the MIS level — I'm not sure that's certain yet; let's just call it a target. Model Office — I think that's relatively certain.

CoT doesn't consume resources. Doing any kind of exploratory research — it doesn't burn through GPUs; it only needs a small number of GPUs. What it needs is ideas. It also doesn't consume talent resources, because you don't need someone constantly stationed on it. It's not a project — what you need is many people thinking about the problem at the same time.

So you don't need to allocate specific resources to it, because it doesn't need resources. Training models, shipping models, running model efficiency experiments — those require resources. The activities I just described — the drain on both people and GPUs — is very small.

We call this “scratch-card hunting.” The barrier is very low — anyone can go scratch. But who scratches up what — I don't know whether that depends on talent or something else.

So this doesn't require us to allocate resources. The only difference between us and other companies is that we spend time discussing this problem, we think about it, and we treat it as something important. Within the company, it's an important problem — one we spend time thinking about. But it doesn't require spending a lot of resources on it.

Next, I also saw a question: the hallucination problem significantly impacts user experience. The hallucination problem can also be solved — there are methods. But it's a long topic. Hallucination, you could say, can be addressed through better post-training — it's a solvable, improvable issue. It's just that no one has put huge effort into it yet.

Or rather, for me, hallucination is a problem, but we categorize it as a product problem. We'll address it — but it's not a top-priority issue.

There was also a question earlier about data labeling. On data labeling — this ties into our capital investment structure. With our capital investment structure, we can't afford the cost of that much high-quality data labeling, because it's very expensive. The cost of data labeling in the US and in China is basically the same. China doesn't have a cost advantage in data labeling — especially in labeling high-end data — which makes it very hard for us to invest in labeling data at the same level as the US.

This path is very difficult in China, because labeling data is just too expensive — whether we outsource it or do it ourselves, it's painful.

LIANG WENFENG · CONTINUED

FOUNDER

So right now, we're basically walking on two legs. It's not that we absolutely can't label — it's that some data is cheaper to label and some is more expensive. We label the cheaper stuff first. So you could say: right now, roughly half of our company is labeling data. Half of our core researchers — the most important people — are involved in data labeling. We focus on labeling data.

Solving the AI problem, at this current stage, depends on data labeling. Just look — it's all about the data.

CHAPTER 05

Products, data and commercialization

User experience, post-training, vertical agents, capital markets and the company's evolving structure.

INVESTOR

INVESTOR

Mr. Liang, thank you. You express things so well — we really feel it deeply. It's excellent. Thank you.

Let me ask a few questions. First: you mentioned that Chinese models are definitely stronger than US models in terms of efficiency. You also mentioned some other areas where China might surpass the US in the future. What do you think those areas are — where we might surpass the US in intelligence or other dimensions?

LIANG WENFENG

FOUNDER

I think in many aspects of user experience, we might be able to do better than the US. I won't speak for others' experiences — our own product experience feels quite good. I think in user experience, we may not be worse than the US. In product, our product capabilities may not be worse than the US. Costs will probably also be lower than the US, so China could very well remain competitive.

In other areas — if there's a structural advantage, I think maybe there isn't one. But in cost and product, I do think there are certain structural advantages. Cost is easy to understand — they don't need to do it, so they don't develop that capability. They definitely don't prioritize it the way we do. We treat it as extremely important; for them, it's not important.

Product too — many companies here have solid product capabilities. So I think these two areas may represent structural advantages.

INVESTOR

INVESTOR

OK. Second question: you mentioned post-training — we invest at relatively high cost, and Anthropic and OpenAI both invest enormous sums. After this funding round, do you think we'll increase investment in post-training?

LIANG WENFENG

FOUNDER

The gap is mainly in high-quality data labeling, and mainly within AI research. We'll definitely increase investment, but high-quality data labeling is not typically a capital-investment problem. The bottleneck for high-quality data labeling, I think, is time — it just takes time.

Because for OpenAI, for the overseas players, for Anthropic — they all started earlier, have more capital, and have more GPUs. Under these conditions, domestically, we can be seen as having only started in the last six months. So in terms of time, I think we need more time.

This isn't strongly tied to capital investment — because even without more capital, the existing capital was already enough to expand at maximum speed. But there's a ceiling on that speed. The bottleneck isn't “I can immediately have more people” — it's not a money bottleneck, and it's not a GPU bottleneck. But it's definitely in a rapid-expansion phase.

So we believe that within a year, the high-quality data problem will be done fairly well — domestically, I think that's a reasonable expectation. The gap may not be that stark, but it really does take time.

INVESTOR

INVESTOR

Thank you. Third question: we've seen Anthropic using their own models for their own products — launching many vertical offerings in finance, legal... and even aiming toward healthcare in the future. Do you think at some stage we will also consider these vertical applications?

LIANG WENFENG

FOUNDER

I haven't thought this through very clearly yet. What our domestic business model will look like in the future — or what the smoothest path will be —

we're not at that stage yet. The domestic situation may not be the same as overseas. What it'll look like here — I think it's still hard to judge right now.

Given the current domestic circumstances, I think the most sensible approach is to go all-in on a general-purpose Agent. Other Agents should have lower priority — including finance, medical, those kinds of Agents. We need to do Coding first, because the Coding Agent can accomplish a lot and covers many vertical Agent use cases. At this stage, we think the most important thing should still be the Coding Agent.

INVESTOR

INVESTOR

Very clear. Thank you.

One more question I wanted to ask. We do DeepSeek — and we also deeply admire you for consistently approaching DeepSeek in a very pure, research-driven way. But now this industry has indeed reached the capital markets stage — it's on the path to capitalization. And you are a very responsible person — whether to your colleagues or to investors, you are very conscientious.

So between the pure research, pure AGI direction and the capital markets — you will certainly need to engage with capital markets in the future, and you'll have public shareholders. How do you see balancing that going forward? How do you think about this?

LIANG WENFENG

FOUNDER

Right now, I think it should be doable — we can have both. If this year I can have a few hundred million dollars in enterprise revenue, plus we have consumer users, then there's already a certain commercial foundation. Next year, with enterprise revenue — if demand can grow further — the company may not be far from net profitability. It might already be at net profitability. It might no longer be a pure cash-burn phase. So I feel there's still quite a lot of room for what can be done, what moves can be made going forward.

Or put it this way: in the worst case, just selling API access could probably support a publicly listed company. If technology doesn't progress further —

if our technology just freezes here — then we go all-in on selling API access, do a good job on those services, and I think that's enough.

So I think there's reason to be confident. It genuinely isn't that hard. Because we're in a high-leverage spot, in a domain developing at breakneck speed — it probably just isn't that hard.

All I can say is: we hope for a bigger dream, but we also have a safety net we can fall back on in terms of performance.

INVESTOR

INVESTOR

Thank you. Very well said. That's all my questions. Thank you.

Thank you for sharing. On some of the topics raised earlier, I want to follow up with you. I think what really sets DeepSeek apart from other companies is that our organization is different — our organizational form is different. But organizational form is related to our goals, and also needs to account for the organization's own boundaries and efficiency.

I wonder, from a macro perspective, is there a good model for our organizational form that we could learn from? Historically, maybe Bell Labs, or some other form might be close to ideal? Or do we ourselves feel there isn't an ideal model, and it's more something we have to gradually explore? Perhaps in a new era, we can only rely on ourselves to explore it.

Because our organizational form is definitely different from any of the US companies, right? The three US companies are themselves different from each other, but at least they all set out from a commercial-company perspective to explore. I mainly want to ask more about the organization question.

LIANG WENFENG

FOUNDER

First, we have no model we're imitating. Every step we take is rooted in our actual situation — seeking truth from facts, making decisions based on real circumstances, and figuring out what we should do. So it's a product of its era, or a reflection of the real conditions. It's not the result of imitation.

What I mean is: under the circumstances, the optimal solution probably was what we chose. Or at least, the path we selected is the one we believe in. Every step, we definitely thought it through, weighed the options, and the outcome is the result of that deliberation. It's not because we saw someone else choose it and then copied them. It's because we analyzed the pros and cons and chose accordingly.

In the future, it'll probably be the same — we're not imitating anyone.

I think we're still different from Bell Labs, because Bell Labs explicitly didn't need commercialization — it was a... But we explicitly need commercialization. At the end of the day, we have to survive. We are, after all, a company — the government won't give us a penny. So we can have a very lofty mission, but ultimately we're a company, and we have to think about how to stay alive.

So the enterprise side is definitely important to us — because in the future, we may have to rely on it to survive. It's just that right now it's not important, because currently it's a cost center. So I think this is still different — different from Bell Labs.

We are fundamentally still a company. Throughout history, there have been many companies that had pursuits beyond profit — but that doesn't mean they weren't companies. Many companies have achieved greatness precisely because they had a pursuit beyond profit. That pursuit, in the end, didn't hurt their commercialization — it actually made their commercialization better.

We are fundamentally still a company. It's just that when it comes to which profits to pursue, when to pursue them, how much to earn, and what to earn them from — we make choices and trade-offs.

CHAPTER 06

TileLang and closing discussion

The final technical exchange on software portability and efficiency, followed by closing remarks.

INVESTOR

INVESTOR

Mr. Liang, I have two quick questions for you. First: you mentioned earlier that MILOS might not be a highly certain target, but you'll definitely head in that direction. And you also mentioned that the activation parameters might be... for the next generation, on the order of 150 to 250B. In that scenario, do you see 150 to 250B benchmarking against O4.7, or against something else? That's the first.

Second question: you also mentioned earlier that on the inference side, you're now using some compilation languages beyond CUDA. My understanding is that originally, based on NVIDIA's ecosystem, you'd mostly use things like PTX. Now that you're using something like TileLang, does that significantly reduce inference efficiency — or at least reduce it in the short term? I'm wondering how you view the efficiency loss from changing compilation languages — or is this actually, in the long run, a net gain that improves efficiency?

LIANG WENFENG

FOUNDER

It improves efficiency — significantly improves efficiency.

INVESTOR

INVESTOR

OK — so there's no negative impact at all?

LIANG WENFENG

FOUNDER

Right — it dramatically improves efficiency. So this is an opportunity. It's like: before, you were locked into the CUDA ecosystem and couldn't leave. Now we can abandon that ecosystem and use a simpler approach — with TileLang. It's a high-level language, and writing those programs is very fast. The amount of code you need to write is very small. I can rewrite the whole thing.

INVESTOR

INVESTOR

Understood. So both of these are actually — as you mentioned — major opportunities brought by AI, not short-term disadvantages that need patching up.

LIANG WENFENG

FOUNDER

Right. It's a major opportunity driven by technology development. It's not just an AI opportunity — we also have a project where we're using AI to write TileLang.

INVESTOR

INVESTOR

Doesn't that make it even faster?

LIANG WENFENG

FOUNDER

Right now, all TileLang is written by humans, but even that is already much faster than writing CUDA before.

INVESTOR

INVESTOR

Understood. So even at the hardware-low-level execution efficiency, you think there's no impact?

LIANG WENFENG

FOUNDER

A 1% to 2% loss — I think that's acceptable.

INVESTOR

INVESTOR

OK. Understood, got it. Alright, thank you. Very clear.

HOST

HOST

Anyone else have questions? If not, let's wrap it up for today.

INVESTOR

INVESTOR

OK, thank you. Thank you everyone for your time. Thank you, Mr. Liang. Thank you.

HOST

HOST

Bye-bye.

LIANG WENFENG

FOUNDER

Bye-bye.

Glossary

Compact definitions for recurring technical terms. These definitions are editorial additions, not part of the transcript.

AGI	Artificial general intelligence: a debated term for systems able to perform a broad range of cognitive tasks at or beyond human level.
AGENT	A system that can plan and carry out multi-step work with tools, feedback and changing intermediate state.
CHAIN OF THOUGHT	Intermediate reasoning used by a model while working toward an answer. Visible explanations and internal reasoning are not always the same.
CONTINUAL LEARNING	The ability to incorporate new experience over time without retraining from scratch or forgetting earlier capabilities.
POST-TRAINING	Training performed after the main pretraining stage to improve instruction following, reasoning, safety and task performance.
MULTIMODALITY	The ability to process or generate more than one type of information, such as text, images, audio or video.
CUDA	NVIDIA's software platform and programming ecosystem for general-purpose computation on its GPUs.
TILELANG	A higher-level language discussed in the transcript as a way to implement efficient GPU kernels with less vendor-specific code.

TRANSPARENCY

Source and editorial status

This is an unofficial public transcript. Its wording and authenticity have not been confirmed by Liang Wenfeng or DeepSeek. The source file labels the opening overview as lightly edited.

TEXT PRESERVATION

This edition preserves the complete user-provided transcript. Chapter titles, glossary definitions and navigation notes are editorial additions and are not part of the meeting record.

ORIGINAL PUBLIC LOCATION

<https://watermelonwater.tech/archives/>

[DeepSeek%20Founder%20Liang%20Wenfeng%3A%20Complete%20Four-Hour%20Investor%20Meeting%20Transcript](https://watermelonwater.tech/archives/DeepSeek%20Founder%20Liang%20Wenfeng%3A%20Complete%20Four-Hour%20Investor%20Meeting%20Transcript)

ATTACHMENT TITLE

DeepSeek Founder Liang Wenfeng: Complete Four-Hour Investor Meeting Transcript

EDITION NOTE

Reader edition typeset from the complete text supplied directly by the user. Typographic dashes and quotation marks were restored for print-quality rendering; no wording was summarized or removed from the transcript body.